

A Real-Time, Flexible Logging Infrastructure for MonPoly

Bachelor's Thesis

Jonas Degelo

ETH Zürich

February 24, 2023

Agenda

Background and Motivation

- Monitors and Databases

- MonPoly

- Time-Series Databases (QuestDB)

The Wrapper

- Architecture

- The Wrapper - Data flow

Algorithm and Policy Change

- Relative Intervals

- Extended Relative Intervals

Partial Policy Change (Work in Progress)

Evaluation

- Overhead of the Wrapper

- Policy Change Optimization

Outlook

Online and Offline Monitoring

- ▶ **Online Monitoring** is done while the system is running
- ▶ **Offline Monitoring** is done on a log after a system has completed
- ▶ We focus on the Online Monitoring problem

Time-Series Databases

- ▶ A time-series database is optimized for the insertion and retrieval of temporal data.
- ▶ QuestDB
 - ▶ Largely compatible with PostgreSQL
 - ▶ Adds a **TIMESTAMP** datatype in addition to the usual types INT, STRING, etc.
- ▶ Dedicated time stamp column (analogous to an index column)



Monitor vs. Database Approach

Monitor (MonPoly)

- ▶ Good when:
 - ▶ Data changes often
 - ▶ Policy changes infrequently
- ▶ Bad when:
 - ▶ Policy changes frequently

Database (QuestDB)

- ▶ Good when:
 - ▶ Data stays almost constant
 - ▶ Varying, frequent queries
- ▶ Bad when:
 - ▶ Data changes rapidly

Motivation

- ▶ Exploit the combination of monitoring and time-series databases
 - ▶ Allows us to offer a policy change method
- ▶ More accessible interface
 - ▶ Web services
 - ▶ Distributed systems
- ▶ Increase fault tolerance

Signature to Database Schema

schema.sql

```
CREATE TABLE P(x1 INT, x2 STRING,  
               tp INT,  
               ts TIMESTAMP)  
timestamp(ts);
```

```
CREATE TABLE Q(x1 INT, x2 INT,  
               tp INT,  
               ts TIMESTAMP)  
timestamp(ts);
```

```
CREATE TABLE R(x1 INT,x2 STRING,  
               x3 STRING,  
               tp INT,  
               ts TIMESTAMP)  
timestamp(ts);
```

```
CREATE TABLE time_points(tp INT,  
                          ts TIMESTAMP)  
timestamp(ts);
```

signature.sig

P(id:int, action:string)

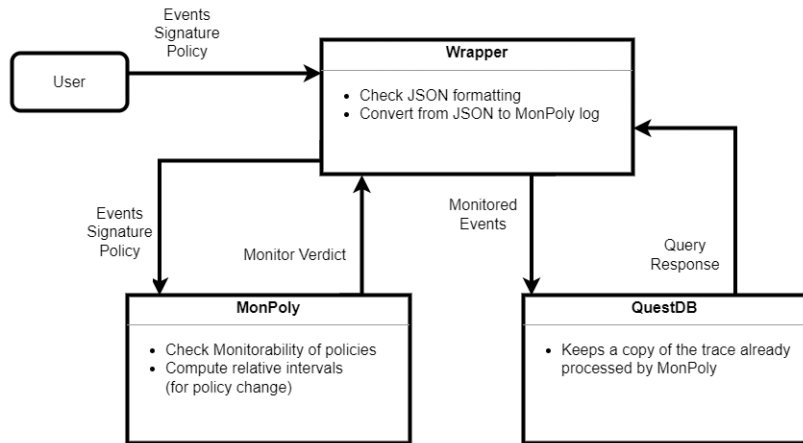
Q(id1:int, id2:int)

R(int, string, string)

The Wrapper - REST API

- | | |
|---------------------|-------------------------|
| ▶ / (Info-Page) | ▶ /db-set-user |
| ▶ /set-policy | ▶ /db-set-password |
| ▶ /set-signature | ▶ /db-set-host |
| ▶ /start-monitor | ▶ /db-set-pgsql-port |
| ▶ /stop-monitor | ▶ /db-set-influxdb-port |
| ▶ /change-policy | ▶ /db-set-database |
| ▶ /get-policy | ▶ /db-get-user |
| ▶ /get-signature | ▶ /db-get-password |
| ▶ /reset-everything | ▶ /db-get-host |
| ▶ /log-events | ▶ /db-get-pgsql-port |
| ▶ /get-events | ▶ /db-get-influxdb-port |
| ▶ /get-most-recent | ▶ /db-get-database |

The Wrapper - Architecture



Policy Change

- ▶ Start a new monitor with the new policy
- ▶ **Goal:** Our monitor evaluates the new policy at the current time point just as if it had seen the same trace as the old monitor
- ▶ **Naive approach:** Read entire trace again
- ▶ **Idea:** Reduce the size of the trace by removing events that do not influence how the new policy gets evaluated

Optimization Example

$$\phi = (P(x, "a") \mathcal{S}_{[0,30)} Q(x, 2)) \wedge (\neg(R(y, "c", z)) \mathcal{S}_{[20,60)} P(y, z))$$

- ▶ **Relative Interval:** $\text{RI}(\phi) = (-60, 0]$
- ▶ **Extended Relative Intervals:** $\text{ERI}(\phi) =$

$$P(*, *) \rightarrow (-60, -20]$$

$$P(*, "a") \rightarrow (-30, 0]$$

$$Q(*, 2) \rightarrow (-30, 0]$$

$$R(*, "c", *) \rightarrow (-60, 0]$$

- ▶ All time points with time stamps not in the interval $\{\text{current time stamp}\} \oplus \text{RI}(\phi)$ do not change the evaluation of ϕ at the current time point.

Interval Operators

Let I and J be two intervals, then

- ▶ $I \oplus J = \{i + j \mid i \in I, j \in J\}$
 - ▶ $[0, 3] \oplus [-2, 4] = [-2, 7]$
- ▶ $I \uplus J$ is the smallest interval that contains all elements that are in at least one of the intervals I and J .
 - ▶ $[-4, 1] \uplus [4, 5] = [-4, 5]$

Relative Intervals

Definition

The relative interval of the formula ϕ , $\text{RI}(\phi) \subseteq \mathbb{Z}$ is defined inductively over the formula structure: $\text{RI}(\phi) =$

$\left\{ \begin{array}{l} \{0\} \\ \text{RI}(\psi) \\ \text{RI}(\psi) \uplus \text{RI}(\chi) \\ (-b, 0] \uplus ((-b, -a] \oplus \text{RI}(\psi)) \\ [0, b) \uplus ([a, b) \oplus \text{RI}(\psi)) \\ (-b, 0] \uplus ((-b, 0] \oplus \text{RI}(\psi)) \uplus ((-b, -a] \oplus \text{RI}(\chi)) \\ [0, b) \uplus ([0, b) \oplus \text{RI}(\psi)) \uplus ([a, b) \oplus \text{RI}(\chi)) \end{array} \right.$	$\begin{array}{l} \text{atomic formulas} \\ \neg\psi, \exists x.\psi, \forall x.\psi \\ \psi \vee \chi, \psi \wedge \chi \\ \bullet_{[a,b)}\psi \\ \bigcirc_{[a,b)} \\ \psi \mathcal{S}_{[a,b)}\chi \\ \psi \mathcal{U}_{[a,b)}\chi \end{array}$
--	--

Basin et al. "Scalable Offline Monitoring of Temporal Specifications". In:
Formal Methods and System Design 49 (1 2016), pp. 75-108

Extended Relative Intervals

Definition

Let M and N be two masked predicate maps and T a positive interval, then

$$\begin{aligned}M \dot{\cup} N &= \{p(I) \rightarrow (I \dot{\cup} J) \mid p(I) \rightarrow I \in m \text{ and } p(I) \rightarrow J \in n\} \\&\quad \cup \{p(I) \rightarrow I \mid (p(I) \rightarrow I \in m \text{ and } p(I) \in k(M) \setminus k(N))\} \\&\quad \cup \{p(I) \rightarrow I \mid (p(I) \rightarrow I \in n \text{ and } p(I) \in k(N) \setminus k(M))\} \\T \dot{\cup} M &= \{p(I) \rightarrow (T \dot{\cup} I) \mid p(I) \rightarrow I \in M\} \\T \dot{\oplus} M &= \{p(I) \rightarrow (T \dot{\oplus} I) \mid p(I) \rightarrow I \in M\}\end{aligned}$$

Relative Intervals

Definition

The relative interval of the formula ϕ , $\text{RI}(\phi) \subseteq \mathbb{Z}$ is defined inductively over the formula structure: $\text{RI}(\phi) =$

$\left\{ \begin{array}{l} \{0\} \\ \text{RI}(\psi) \\ \text{RI}(\psi) \uplus \text{RI}(\chi) \\ (-b, 0] \uplus ((-b, -a] \oplus \text{RI}(\psi)) \\ [0, b) \uplus ([a, b) \oplus \text{RI}(\psi)) \\ (-b, 0] \uplus ((-b, 0] \oplus \text{RI}(\psi)) \uplus ((-b, -a] \oplus \text{RI}(\chi)) \\ [0, b) \uplus ([0, b) \oplus \text{RI}(\psi)) \uplus ([a, b) \oplus \text{RI}(\chi)) \end{array} \right.$	$\begin{array}{l} \text{atomic formulas} \\ \neg\psi, \exists x.\psi, \forall x.\psi \\ \psi \vee \chi, \psi \wedge \chi \\ \bullet_{[a,b)}\psi \\ \bigcirc_{[a,b)} \\ \psi \mathcal{S}_{[a,b)}\chi \\ \psi \mathcal{U}_{[a,b)}\chi \end{array}$
--	--

Basin et al. "Scalable Offline Monitoring of Temporal Specifications". In:
Formal Methods and System Design 49 (1 2016), pp. 75-108

Extended Relative Intervals

ERI(ϕ) =

$$\left\{ \begin{array}{ll} \{\} & \text{atomic formulas} \\ \{p(m) \rightarrow [0, 0]\} & \text{excl. predicates} \\ \text{ERI}(\psi) & \text{predicate } p \\ & \text{with mask } m, \\ & \neg\psi, \exists x.\psi, \forall x.\psi \\ \text{ERI}(\psi) \dot{\cup} \text{ERI}(\chi) & \psi \vee \chi, \psi \wedge \chi \\ (-b, 0] \dot{\cup} ((-b, -a] \dot{\oplus} \text{ERI}(\psi)) & \bullet_{[a,b]}\psi \\ [0, b) \dot{\cup} ([a, b) \dot{\oplus} \text{ERI}(\psi)) & \circ_{[a,b]}\psi \\ (-b, 0] \dot{\cup} ((-b, 0] \dot{\oplus} \text{ERI}(\psi)) \dot{\cup} ((-b, -a] \dot{\oplus} \text{ERI}(\chi)) & \psi \mathcal{S}_{[a,b]} \chi \\ [0, b) \dot{\cup} ([0, b) \dot{\oplus} \text{ERI}(\psi)) \dot{\cup} ([a, b) \dot{\oplus} \text{ERI}(\chi)) & \psi \mathcal{U}_{[a,b]} \chi \end{array} \right.$$

We proofed that this definition is a correct approximation such that it extracts all necessary time points from a trace and that the evaluation of ϕ is the same for the original trace and the extraction.

Extended Relative Intervals - SQL Query

$$\phi = (P(x, "a") \mathcal{S}_{[0,30)} Q(x, 2)) \wedge (\neg(R(y, "c", z)) \mathcal{S}_{[20,60)} P(y, z))$$

- ▶ **Relative Interval:** $\text{RI}(\phi) = (-60, 0]$
- ▶ **Extended Relative Intervals:** $\text{ERI}(\phi) =$

$$P(*, *) \rightarrow (-60, -20]$$

$$P(*, "a") \rightarrow (-30, 0]$$

$$Q(*, 2) \rightarrow (-30, 0]$$

$$R(*, "c", *) \rightarrow (-60, 0]$$

Extended Relative Intervals - SQL Query

```
SELECT * FROM P WHERE
    (time_stamp > <NOW>-60 AND time_stamp <= <NOW>-20)
OR
    (x2 = "a" AND
     time_stamp > <NOW> -30 AND time_stamp <= <NOW>-0);
SELECT * FROM Q WHERE
    (x2 = 2 AND
     time_stamp > <NOW>-30 AND time_stamp <= <NOW>-0);
SELECT * FROM R WHERE
    (x2 = "c" AND
     time_stamp > <NOW>-60 AND time_stamp <= <NOW>-0);
SELECT * FROM time_points WHERE
    (time_stamp > <NOW>-60 AND time_stamp <= <NOW>-0);
```

Database Result Conversion

QuestDB Response

```
[ "R": [],  
  "Q": [[0,8,65536,"Thu, 01 Jan 1970 00:00:01 GMT"]],  
  "P": [],  
  "time_points": [[65536,"Thu, 01 Jan 1970 00:00:01 GMT"],  
                   [65537,"Thu, 01 Jan 1970 00:00:02 GMT"]]
```

wrapper.json

```
[ "predicates": [ "name": "Q", "occurrences": [[0,8]] ],  
  "timepoint": 65536, "timestamp": "1970-01-01 00:00:01",  
  "predicates": [ ], "timepoint": 65537, "timestamp": "1970-01-01  
00:00:02"]
```

monpoly.log

```
@1 Q(0, 8) ;  
@2 ;
```

Partial Policy Change in MonPoly (Work in Progress)

- ▶ Only for **First Order Logic** operators above temporal operators
- ▶ **Named formulas:**
NAME[f1, name1] OR NAME[f2 and f3, name2]
- ▶ Added *data types* for NAME in MonPoly
- ▶ Updated *formula parser* for NAME constructs
- ▶ Started work on commands for adding and removing parts of formulas.
- ▶ **Commands** to add or remove conjuncts or disjuncts:

```
>remove_part <name><
```

```
>add_conjunct <name1> <name2> <formula><
```

```
>add_disjunct <name1> <name2> <formula><
```

Partial Policy Change in MonPoly (Work in Progress)

$$\phi = (P(x, "a") \mathcal{S}_{[0,30)} Q(x, 2)) \wedge (\neg(R(y, "c", z)) \mathcal{S}_{[20,60)} P(y, z))$$

Could become:

NAME[(P(x,"a") SINCE[0,30) Q(x,2)), **part1**]

AND

NAME[(NOT (R(y,"c",z)) SINCE[20,60) P(y,z)), **part2**]

>add_disjunct part1 part3 'NOT P(2,"b")'<

(**NAME**[P(x,"a") SINCE[0,30) Q(x,2), **part1**]

OR

NAME[NOT P(2,"b"), **part3**])

AND

NAME[NOT (R(y,"c",z)) SINCE[20,60) P(y,z), **part2**]

Partial Policy Change in MonPoly (Work in Progress)

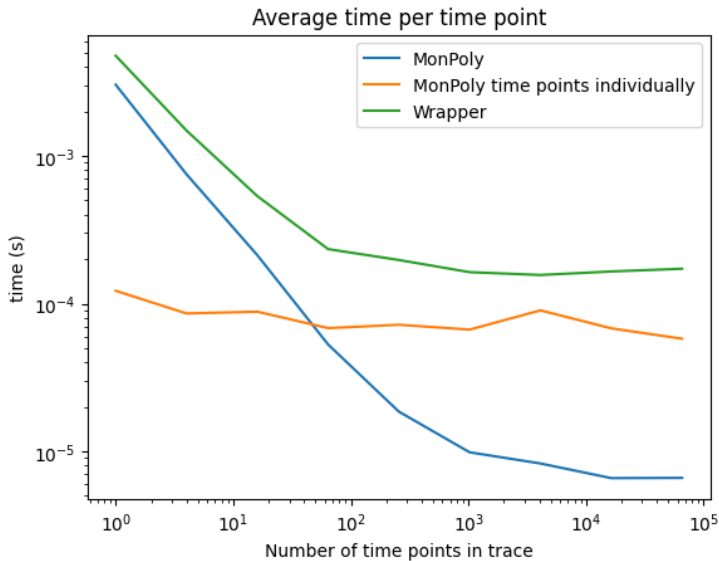
Next up:

- ▶ Compute the internal state for formula parts that will be added.
- ▶ Combine existing state with the state of the new formula.
- ▶ Update state when a formula part gets removed

Overhead of the Wrapper

- ▶ Random policy with intervals bounded by $[0, 20)$
- ▶ Random trace with 4^i time points, for $i = 0, \dots, 8$
- ▶ Measure the time it takes for MonPoly to monitor these traces
- ▶ Measure the time it takes for MonPoly to monitor the traces, when the time points are not in a single file, but sent individually
 - ▶ This is a better simulation of online monitoring
- ▶ Measure the time it takes for the wrapper to monitor these traces
- ▶ This was done **5 times** (with 5 different policies)

Overhead of the Wrapper




Overhead of the Wrapper - Conclusion

- ▶ More detailed profiling is required
- ▶ Possible sources of the overhead are:
 - ▶ The conversion of our JSON format for logs to the MonPoly format
 - ▶ The wrapper sends time points synchronously to MonPoly
 - ▶ The wrapper must go over the events once to send them to MonPoly and again later to send them to QuestDB

Policy Change Optimization

- ▶ Policy change with 2 randomly generated formulas
- ▶ For every pair of formulas the policy change is done with random traces of different length
- ▶ For the policies all temporal operators have intervals that are at most $[0, 20)$
- ▶ The trace length refers to the number of time points in the trace
- ▶ We used traces with lengths 4^i for $i = 0, \dots, 8$
- ▶ **In conclusion:** 9 times policy change for each of the 5 formula pairs

Policy Change Optimization



Outlook

- ▶ Reduce overhead of the wrapper
 - ▶ More in depth profiling of the wrapper
 - ▶ Send time points asynchronously (don't wait for response before sending the next time point)
- ▶ Speed up policy change
 - ▶ More involved checking for constraints on variables in formulas (such as $x < 20$)
- ▶ Continue the partial policy change within MonPoly