

A Real-Time, Flexible Logging Infrastructure for MonPoly

Bachelor's Thesis

Jonas Degelo

ETH Zürich

February 24, 2023

Monitoring Problem

- ▶ **Runtime Monitoring:** Checking a systems actual behavior against a formal specification
- ▶ **Online Monitoring** is done while the system is running
- ▶ **Offline Monitoring** is done on a log after a system has completed
- ▶ We focus on the Online Monitoring problem

Time-Series Database

- ▶ A time-series database is optimized for the insertion and retrieval of temporal data.
- ▶ QuestDB
 - ▶ Largely compatible with PostgreSQL
 - ▶ Adds a **TIMESTAMP** datatype in addition to the usual types INT, STRING, etc.
- ▶ Dedicated time stamp column (analogous to an index column)



Monitor vs. Database Approach

Monitor (MonPoly)

- ▶ Good when:
 - ▶ Data changes often
 - ▶ Policy changes infrequently
- ▶ Bad when:
 - ▶ Policy changes frequently

Time-Series Database (QuestDB)

- ▶ Good when:
 - ▶ Data stays almost constant
 - ▶ Varying, frequent queries
- ▶ Bad when:
 - ▶ Data changes rapidly

Combining Monitoring and Time-Series Database: Motivation

- ▶ Build a wrapper that allows for both logging and monitoring on the same data
- ▶ Provides the functionality of both monitors and time-series database
- ▶ Allows us to offer a policy change method
- ▶ More accessible interface
 - ▶ Web services
 - ▶ Distributed systems
- ▶ Increase fault tolerance

Signature to Database Schema

schema.sql

```
CREATE TABLE P(x1 INT,x2 STRING,  
                tp INT,  
                ts TIMESTAMP)  
    timestamp(ts);  
  
CREATE TABLE Q(x1 INT,x2 INT,  
                tp INT,  
                ts TIMESTAMP)  
    timestamp(ts);  
  
CREATE TABLE R(x1 INT,x2 STRING,  
                x3 STRING,  
                tp INT,  
                ts TIMESTAMP)  
    timestamp(ts);  
  
CREATE TABLE time_points  
    (tp INT,  
     ts TIMESTAMP)  
    timestamp(ts);
```

signature.sig

```
P(id:int,act:string)  
Q(id1:int,id2:int)  
R(int,string,string)
```

Agenda

Background and Motivation

- MonPoly

- Time-Series Database (QuestDB)

The Wrapper

- Architecture

- Data flow

Algorithm and Policy Change

- Relative Intervals

- Extended Relative Intervals

Partial Policy Change (Work in Progress)

Evaluation

- Overhead of the Wrapper

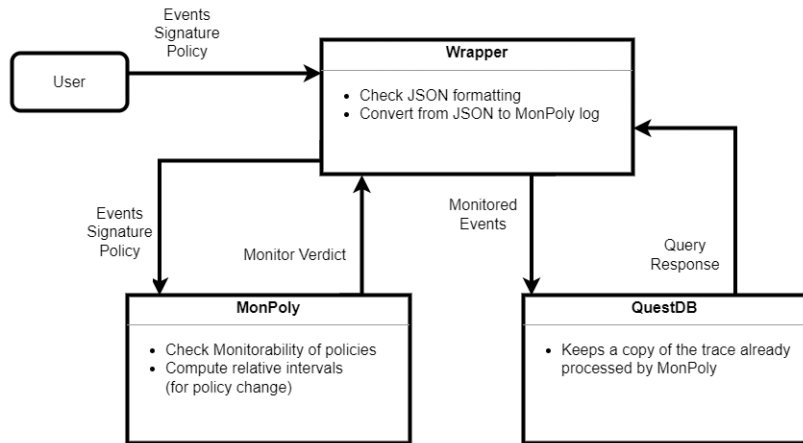
- Policy Change Optimization

Outlook

The Wrapper - REST API

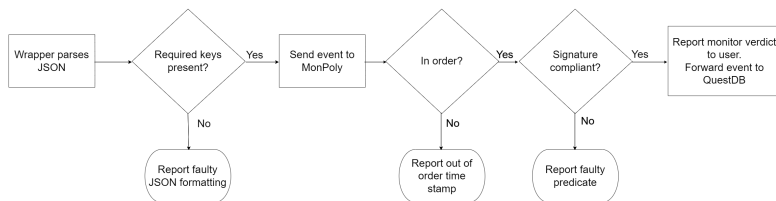
- ▶ / (Info-Page)
- ▶ /set-policy
- ▶ /set-signature
- ▶ /start-monitor
- ▶ /stop-monitor
- ▶ /change-policy
- ▶ /get-policy
- ▶ /get-signature
- ▶ /reset-everything
- ▶ /log-events
- ▶ /get-events
- ▶ /get-most-recent
- ▶ /db-set-user
- ▶ /db-set-password
- ▶ /db-set-host
- ▶ /db-set-pgsql-port
- ▶ /db-set-influxdb-port
- ▶ /db-set-database
- ▶ /db-get-user
- ▶ /db-get-password
- ▶ /db-get-host
- ▶ /db-get-pgsql-port
- ▶ /db-get-influxdb-port
- ▶ /db-get-database

The Wrapper - Architecture



Data flow

- ▶ Increased fault tolerance (compared to MonPoly)



Policy Change

- ▶ Start a new monitor with the new policy
- ▶ **Goal:** Our monitor evaluates the new policy at the current time point just as if it had seen the same trace as the old monitor
- ▶ **Naive approach:** Read entire trace again
- ▶ **Idea:** Reduce the size of the trace by removing events that do not influence how the new policy gets evaluated

Optimization Example

$$\phi = (P(x, "a") \mathcal{S}_{[0,30)} Q(x, 2)) \wedge (\neg(R(y, "c", z)) \mathcal{S}_{[20,60)} P(y, z))$$

- ▶ **Relative Interval:** $\text{RI}(\phi) = (-60, 0]$
- ▶ **Extended Relative Intervals:** $\text{ERI}(\phi) =$

$$P(*, *) \rightarrow (-60, -20]$$

$$P(*, "a") \rightarrow (-30, 0]$$

$$Q(*, 2) \rightarrow (-30, 0]$$

$$R(*, "c", *) \rightarrow (-60, 0]$$

- ▶ All time points with time stamps not in the interval $\{\text{current time stamp}\} \oplus \text{RI}(\phi)$ do not change the evaluation of ϕ at the current time point.

Interval Operators

Let I and J be two intervals, then

- ▶ $I \oplus J = \{i + j \mid i \in I, j \in J\}$
 - ▶ $[0, 3] \oplus [-2, 4] = [-2, 7]$
- ▶ $I \uplus J$ is the smallest interval that contains all elements that are in at least one of the intervals I and J .
 - ▶ $[-4, 1] \uplus [4, 5] = [-4, 5]$

Relative Intervals

Definition

The relative interval of the formula ϕ , $\text{RI}(\phi) \subseteq \mathbb{Z}$ is defined inductively over the formula structure: $\text{RI}(\phi) =$

$\left\{ \begin{array}{l} \{0\} \\ \text{RI}(\psi) \\ \text{RI}(\psi) \uplus \text{RI}(\chi) \\ (-b, 0] \uplus ((-b, -a] \oplus \text{RI}(\psi)) \\ [0, b) \uplus ([a, b) \oplus \text{RI}(\psi)) \\ (-b, 0] \uplus ((-b, 0] \oplus \text{RI}(\psi)) \uplus ((-b, -a] \oplus \text{RI}(\chi)) \\ [0, b) \uplus ([0, b) \oplus \text{RI}(\psi)) \uplus ([a, b) \oplus \text{RI}(\chi)) \end{array} \right.$	$\begin{array}{l} \text{atomic formulas} \\ \neg\psi, \exists x.\psi, \forall x.\psi \\ \psi \vee \chi, \psi \wedge \chi \\ \bullet_{[a,b)}\psi \\ \bigcirc_{[a,b)} \\ \psi \mathcal{S}_{[a,b)}\chi \\ \psi \mathcal{U}_{[a,b)}\chi \end{array}$
--	--

Basin et al. "Scalable Offline Monitoring of Temporal Specifications". In:
Formal Methods and System Design 49 (1 2016), pp. 75-108

Extended Relative Intervals

Map data structure replaces intervals and new operators are needed. Let M and N be two masked predicate maps and T an interval, then

$$\begin{aligned}M \dot{\cup} N &= \{p(I) \rightarrow (I \dot{\cup} J) \mid p(I) \rightarrow I \in m \text{ and } p(I) \rightarrow J \in n\} \\&\quad \cup \{p(I) \rightarrow I \mid (p(I) \rightarrow I \in m \text{ and } p(I) \in k(M) \setminus k(N))\} \\&\quad \cup \{p(I) \rightarrow I \mid (p(I) \rightarrow I \in n \text{ and } p(I) \in k(N) \setminus k(M))\} \\T \dot{\cup} M &= \{p(I) \rightarrow (T \dot{\cup} I) \mid p(I) \rightarrow I \in M\} \\T \dot{\oplus} M &= \{p(I) \rightarrow (T \dot{\oplus} I) \mid p(I) \rightarrow I \in M\}\end{aligned}$$

Relative Intervals

Definition

The relative interval of the formula ϕ , $\text{RI}(\phi) \subseteq \mathbb{Z}$ is defined inductively over the formula structure: $\text{RI}(\phi) =$

$\left\{ \begin{array}{l} \{0\} \\ \text{RI}(\psi) \\ \text{RI}(\psi) \uplus \text{RI}(\chi) \\ (-b, 0] \uplus ((-b, -a] \oplus \text{RI}(\psi)) \\ [0, b) \uplus ([a, b) \oplus \text{RI}(\psi)) \\ (-b, 0] \uplus ((-b, 0] \oplus \text{RI}(\psi)) \uplus ((-b, -a] \oplus \text{RI}(\chi)) \\ [0, b) \uplus ([0, b) \oplus \text{RI}(\psi)) \uplus ([a, b) \oplus \text{RI}(\chi)) \end{array} \right.$	$\begin{array}{l} \text{atomic formulas} \\ \neg\psi, \exists x.\psi, \forall x.\psi \\ \psi \vee \chi, \psi \wedge \chi \\ \bullet_{[a,b)}\psi \\ \bigcirc_{[a,b)} \\ \psi \mathcal{S}_{[a,b)}\chi \\ \psi \mathcal{U}_{[a,b)}\chi \end{array}$
--	--

Basin et al. "Scalable Offline Monitoring of Temporal Specifications". In:
Formal Methods and System Design 49 (1 2016), pp. 75-108

Extended Relative Intervals

$\text{ERI}(\phi) =$

$\{\}$	atomic formulas
$\{p(m) \rightarrow [0, 0]\}$	excl. predicates
$\text{ERI}(\psi)$	predicate p
$\text{ERI}(\psi) \dot{\cup} \text{ERI}(\chi)$	with mask m ,
$(-b, 0] \dot{\cup} ((-b, -a] \dot{\oplus} \text{ERI}(\psi))$	$\neg\psi, \exists x.\psi, \forall x.\psi$
$[0, b) \dot{\cup} ([a, b) \dot{\oplus} \text{ERI}(\psi))$	$\psi \vee \chi, \psi \wedge \chi$
$(-b, 0] \dot{\cup} ((-b, 0] \dot{\oplus} \text{ERI}(\psi)) \dot{\cup} ((-b, -a] \dot{\oplus} \text{ERI}(\chi))$	$\bullet_{[a,b)}\psi$
$[0, b) \dot{\cup} ([0, b) \dot{\oplus} \text{ERI}(\psi)) \dot{\cup} ([a, b) \dot{\oplus} \text{ERI}(\chi))$	$\bigcirc_{[a,b)}\psi$
	$\psi \mathcal{S}_{[a,b)} \chi$
	$\psi \mathcal{U}_{[a,b)} \chi$

We proofed that this definition is a correct approximation such that it extracts all necessary time points from a trace and that the evaluation of ϕ is the same for the original trace and the extraction.

Extended Relative Intervals - SQL Query

$$\phi = (P(x, "a") \mathcal{S}_{[0,30)} Q(x, 2)) \wedge (\neg(R(y, "c", z)) \mathcal{S}_{[20,60)} P(y, z))$$

- ▶ **Relative Interval:** $\text{RI}(\phi) = (-60, 0]$
- ▶ **Extended Relative Intervals:** $\text{ERI}(\phi) =$

$$P(*, *) \rightarrow (-60, -20]$$

$$P(*, "a") \rightarrow (-30, 0]$$

$$Q(*, 2) \rightarrow (-30, 0]$$

$$R(*, "c", *) \rightarrow (-60, 0]$$

Extended Relative Intervals - SQL Query

```
SELECT * FROM P WHERE
    (time_stamp > <NOW>-60 AND time_stamp <= <NOW>-20)
OR
    (x2 = "a" AND
     time_stamp > <NOW> -30 AND time_stamp <= <NOW>-0);
SELECT * FROM Q WHERE
    (x2 = 2 AND
     time_stamp > <NOW>-30 AND time_stamp <= <NOW>-0);
SELECT * FROM R WHERE
    (x2 = "c" AND
     time_stamp > <NOW>-60 AND time_stamp <= <NOW>-0);
SELECT * FROM time_points WHERE
    (time_stamp > <NOW>-60 AND time_stamp <= <NOW>-0);
```

Database Result Conversion

QuestDB Response

```
[{"R": []},  
{"Q": [[0,8,65536,"Thu, 01 Jan 1970 00:00:01 GMT"]]},  
{"P": []},  
{"time_points": [[65536,"Thu, 01 Jan 1970 00:00:01 GMT"],  
[65537,"Thu, 01 Jan 1970 00:00:02 GMT"]]}]
```

wrapper.json

```
[{"predicates": [{"name": "Q", "occurrences": [[0,8]]}],  
"timepoint": 65536, "timestamp": "1970-01-01 00:00:01"},  
{"predicates": [],  
"timepoint": 65537, "timestamp": "1970-01-01 00:00:02"}]
```

monpoly.log

```
@1 Q(0,8);  
@2 ;
```

Partial Policy Change in MonPoly (Work in Progress)

- ▶ Only for **First Order Logic** operators above temporal operators
- ▶ **Named formulas:**
NAME[f1, name1] OR NAME[f2 and f3, name2]
- ▶ Added *data types* for NAME in MonPoly
- ▶ Updated *formula parser* for NAME constructs
- ▶ Started work on commands for adding and removing parts of formulas.
- ▶ **Commands** to add or remove conjuncts or disjuncts:

```
>remove_part <name><
```

```
>add_conjunct <name1> <name2> <formula><
```

```
>add_disjunct <name1> <name2> <formula><
```

Partial Policy Change in MonPoly (Work in Progress)

$$\phi = (P(x, "a") \mathcal{S}_{[0,30)} Q(x, 2)) \wedge (\neg(R(y, "c", z)) \mathcal{S}_{[20,60)} P(y, z))$$

Becomes:

```
NAME[(P(x,"a") SINCE[0,30) Q(x,2)), part1] AND  
NAME[(NOT (R(y,"c",z)) SINCE[20,60) P(y,z)), part2]
```

```
>add_disjunct part1 part3 'NOT P(2,"b")'<
```

```
(NAME[P(x,"a") SINCE[0,30) Q(x,2), part1] OR  
NAME[NOT P(2,"b"), part3]) AND  
NAME[NOT (R(y,"c",z)) SINCE[20,60) P(y,z), part2]
```

$$\phi' = ((P(x, "a") \mathcal{S}_{[0,30)} Q(x, 2)) \wedge (\neg P(2, "b"))) \\ \wedge (\neg(R(y, "c", z)) \mathcal{S}_{[20,60)} P(y, z))$$

Partial Policy Change in MonPoly (Work in Progress)

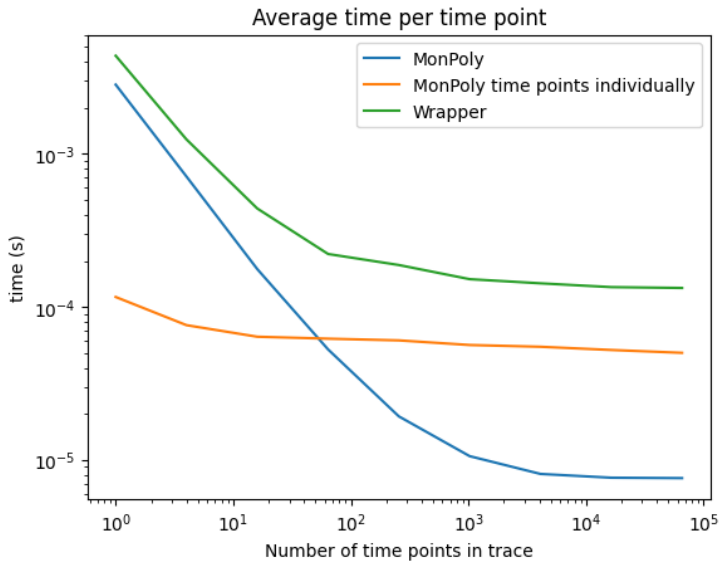
Next up:

- ▶ Compute the internal state for formula parts that will be added.
- ▶ Combine existing state with the state of the new formula.
- ▶ Update state when a formula part gets removed.

Overhead of the Wrapper

- ▶ Random policy with intervals bounded by $[0, 20)$
- ▶ Random trace with 4^i time points, for $i = 0, \dots, 8$
- ▶ Measure the time it takes for MonPoly to monitor these traces
- ▶ Measure the time it takes for MonPoly to monitor the traces, when the time points are not in a single file, but sent individually
 - ▶ This is a better simulation of online monitoring
- ▶ Measure the time it takes for the wrapper to monitor these traces
- ▶ This was done **100 times** (with 100 different policies)

Overhead of the Wrapper



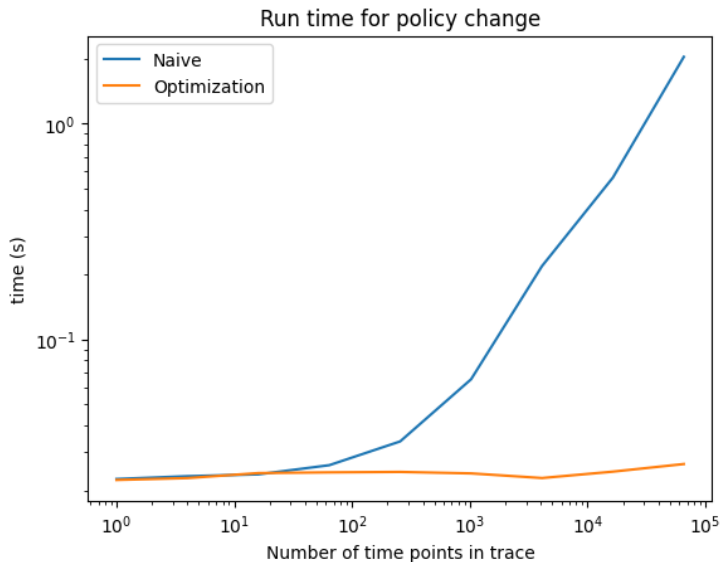
Overhead of the Wrapper - Conclusion

- ▶ More detailed profiling is required
- ▶ Possible sources of the overhead are:
 - ▶ The conversion of our JSON format for logs to the MonPoly format
 - ▶ The wrapper sends time points synchronously to MonPoly
 - ▶ The wrapper must go over the events once to send them to MonPoly and again later to send them to QuestDB

Policy Change Optimization

- ▶ Policy change with 2 randomly generated formulas
- ▶ Intervals in policies are bounded by $[0, 20)$.
- ▶ Random traces with 4^i time points, for $i = 0, \dots, 8$
- ▶ Load trace into wrapper and then initiate the policy change
- ▶ This was done 100 times, for 100 different policy pairs.

Policy Change Optimization



Outlook

- ▶ Reduce overhead of the wrapper
 - ▶ More in depth profiling of the wrapper
 - ▶ Send time points asynchronously (don't wait for response before sending the next time point)
- ▶ Speed up policy change
 - ▶ More involved checking for constraints on variables in formulas (such as $x < 20$)
- ▶ Continue the partial policy change within MonPoly

Thank you for your attention!